

HADOOP MULTINODE INSTALLATION

Multi node-cluster setup

We will create a 3 node cluster system (1-Master, 2-slave).

Systems should be connected to same network

Check if nodes are reachable

Find the ip Address of all 3 systems & try to ping each other

hduser@ubuntu:~\$ifconfig

Master 192.168.220.180

Slave1 192.168.220.181

Slave2 192.168.220.182

Master hduser@ubuntu:~\$ ping 192.168.220.181 //Master pinging slave1

Master hduser@ubuntu:~\$ ping 192.168.220.182 //Master pinging slave2

Slave1 hduser@ubuntu:~\$ ping 192.168.220.180 //Slave1 pinging master

Slave1 hduser@ubuntu:~\$ ping 192.168.220.182 //Slave1 pinging slave1

Slave2 hduser@ubuntu:~\$ ping 192.168.220.180 //Slave2 pinging master

Slave2 hduser@ubuntu:~\$ ping 192.168.220.181 //Slave2 pinging slave1

Change the hostname of all 3 system

Go to Master system.

\$ sudo vim /etc/hostname

Replace ubuntu with master in file.

Press ESC on the keyboard.

Save the configuration by :wq!

Repeat the above step with slave1 & slave2 VM and change the host name to **slave1**, **slave2**

Update the hosts on all 3 nodes

Master SYSTEM:

\$ sudo vim /etc/hosts

127.0.0.1 localhost # dont make any changes to this line.

#127.0.1.1 master # remove this line.

REFERENCE :KAUSHIK SHAKKARI
EMAIL : kaushik.shakkari@gmail.com

HADOOP MULTINODE INSTALLATION

```
192.168.220.185 master      #Add this 3 lines.  
192.168.220.186 slave1  
192.168.220.187 slave2
```

Repeat the same in Slave1 and Slave2

1. Restart all systems

Confirm the hostname of all 3 nodes.

Execute below command in all three systems.

\$hostname

It should print **master, slave1, slave2** in 3 machines respectively.

2. Ping Each other using Hostname

Start pinging each other system again using the hostname instead of ipaddress

Master->ping slave1 & slave2

Slave1->ping master only

Slave2->ping master only

```
hduser@master:~$ ping slave1  
hduser@master:~$ ping slave2  
hduser@slave1 :~$ ping master  
hduser@slave2 :~$ ping master
```

3. Test SSH connectivity

```
hduser@master:~$ ssh master  
hduser@master:~$ ssh slave1  
hduser@slave1:~$ exit  
master.  
hduser@master:~$ ssh slave2  
hduser@slave2:~$ exit  
master.  
hduser@master:~$
```

**#Type Yes. It will connect to slave1
#will exit slave1 and goes back to**

**# Type Yes. It will connect to slave2
#will exit slave2 and goes back to**

HADOOP MULTINODE INSTALLATION

It will ask for yes or no and you should type 'yes'

We should be able to SSH master and SSH slaves without password prompt.

If it asks for password while connecting to master or slave using SSH, there is something went wrong and you need to fix it before proceeding further.

Common Configurations (needs to be done in master as well as all slaves)

1. Update core-site.xml(Master+ All Slave Nodes)

sudo vim /usr/local/hadoop/etc/hadoop/core-site.xml

we have 2 changes

- a. Remove hadoop.tmp.dir configuration. We don't require them
- b. change localhost to master
replace

```
<property>
  <name>hadoop.tmp.dir</name>
  <value>/app/hadoop/tmp</value>
  <description>A base for other temporary directories.</description>
</property>
```

with

```
<property>
  <name> fs.default.name </name>
  <value>hdfs://master:9000</value> <!--instead of localhost -->
</property>
```

2. Update hdfs-site.xml(Master + All slave Nodes)

3 changes

- a. Replication is set to 2
- b. Namenode configured only in master
- c. Datanode configured only in slave

\$ sudo vim /usr/local/hadoop/etc/hadoop/hdfs-site.xml

```
<property>
  <name>dfs.replication</name>
  <value>2</value> <!--changing replication from 1 to 2 -->
</property>
```

<!--Keep this entry in Master only, and delete from slaves-->

```
<property>
  <name>dfs.namenode.name.dir</name>
  <value>file:/usr/local/hadoop_tmp/hdfs/namenode</value>
```

HADOOP MULTINODE INSTALLATION

```
</property>
```

```
<!--Remove this entry From Master, and Keep this entry in slaves-->
```

```
<property>
```

```
  <name>dfs.datanode.data.dir</name>
```

```
  <value>file:/usr/local/hadoop_tmp/hdfs/datanode</value>
```

```
</property>
```

3. Update yarn-site.xml(Master + All Slave Nodes)

```
$ sudo vim /usr/local/hadoop/etc/hadoop/yarn-site.xml
```

```
<property>
```

```
  <name>yarn.resourcemanager.resource-tracker.address</name>
```

```
  <value>master:8025</value>
```

```
</property>
```

```
<property>
```

```
  <name>yarn.resourcemanager.scheduler.address</name>
```

```
  <value>master:8030</value>
```

```
</property>
```

```
<property>
```

```
  <name>yarn.resourcemanager.address</name>
```

```
  <value>master:8050</value>
```

```
</property>
```

4. Update mapred-site.xml(Master + All Slave Nodes)

```
$sudo vim /usr/local/hadoop/etc/hadoop/mapred-site.xml
```

```
  <property>
```

```
    <name>mapreduce.jobhistory.address</name>
```

```
    <value>master:10020</value><!-- Pointing jobhistory to master-->
```

```
    <description>Host and port for Job History Server (default
```

```
    0.0.0.0:10020)</description>
```

```
  </property>
```

Master only Configuration

5. Update Masters and slaves file(Master Node only)

If you see any entry related to localhost, feel free to delete it. This file is just helper file that are used by hadoop scripts to start appropriate services on master and slave nodes.

```
hduser@master$ sudo vim /usr/local/hadoop/etc/hadoop/slaves
```

REFERENCE :KAUSHIK SHAKKARI
EMAIL : kaushik.shakkari@gmail.com

HADOOP MULTINODE INSTALLATION

slave1

slave2

Below masters file does not exists by default. It gets created the file.

```
hduser@master$ sudo vim /usr/local/hadoop/etc/hadoop/masters
master
```

Note: You don't need to configure them in slave nodes

6. Recreate Namenode folder(Master Only)

```
hduser@master$
sudo rm -rf /usr/local/hadoop_tmp
sudo mkdir -p /usr/local/hadoop_tmp/hdfs/namenode
sudo chown hduser:hadoop -R /usr/local/hadoop_tmp/
sudo chmod 777 /usr/local/hadoop_tmp/hdfs/namenode
```

Recreate Datanode folder(All Slave Nodes Only)

```
hduser@slave1$
sudo rm -rf /usr/local/hadoop_tmp
sudo mkdir -p /usr/local/hadoop_tmp/hdfs/datanode
sudo chown hduser:hadoop -R /usr/local/hadoop_tmp/
sudo chmod 777 /usr/local/hadoop_tmp/hdfs/datanode
```

7. Format the Name node(Master only)

Before starting the cluster, we need to format the Name node. Use the following command only on master node:

```
$ hdfs namenode -format
```

8. Start the DFS & Yarn (Master Only)

```
$ start-dfs.sh
$ start-yarn.sh
```

or

```
$ start-all.sh #deprecated but internally calls start-dfs.sh && start-yarn.sh
Type Yes when asked for.
```

You should observe that it tries to start data node on slave nodes one by one.

Once it is started, Do a Jps on Master and slaves.

Jps on Master node

HADOOP MULTINODE INSTALLATION

```
hduser@master$ jps
3379 NameNode           #because of start-dfs.sh
3175 SecondaryNameNode  #because of start-dfs.sh
3539 ResourceManager    #because of start-yarn.sh
```

Jps on slave nodes(slave1 and slave2)

```
hduser@slave1$ jps
2484 DataNode           #because of start-dfs.sh
2607 NodeManager        #because of start-yarn.sh
```

9. Review Yarn console:

If all the services started successfully on all nodes, then you should see all of your nodes listed under Yarn nodes. You can hit the following url on your browser and verify that:

<http://master:8088/cluster/nodes>

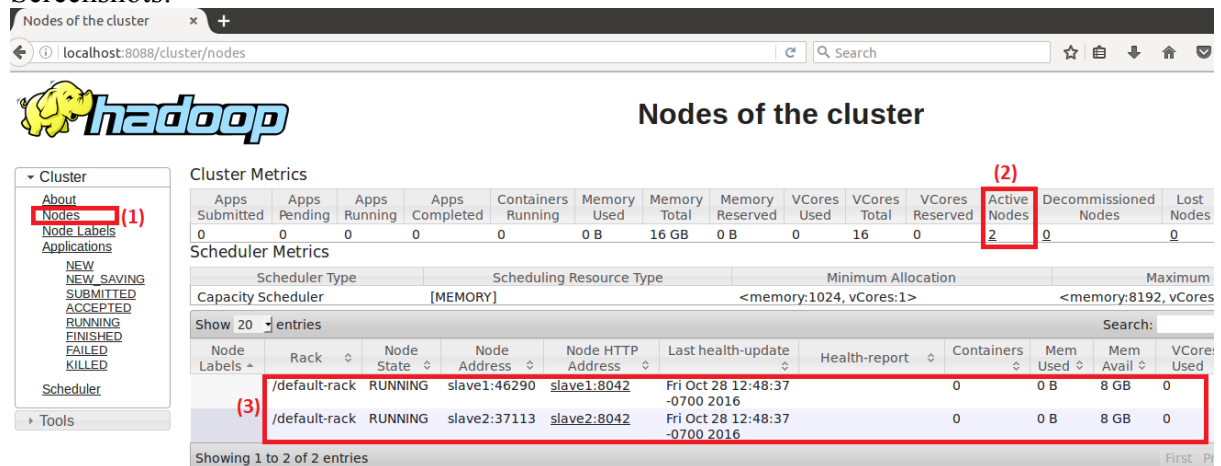
<http://master:50070> # can show live node count and info about each live nodes.

You can also get the report of your cluster by issuing the below commands

```
hduser@master$ hdfs dfsadmin -report
```

In case you don't get to see the live nodes on the browsers, there are some mistakes somewhere. You need to figure it out. You can look into the logs in the below location /usr/local/hadoop/logs/
Check in all the log files of both master and slaves and it should hint you the problem.

Screenshots:



The screenshot shows the Hadoop Yarn console interface. On the left, a sidebar menu has 'Nodes' highlighted with a red box (1). The main content area is titled 'Nodes of the cluster'. It includes a 'Cluster Metrics' table with 'Active Nodes' highlighted by a red box (2). Below this is a 'Scheduler Metrics' section. At the bottom, a table lists active nodes, with two rows highlighted by a red box (3):

Node Labels	Rack	Node State	Node Address	Node HTTP Address	Last health-update	Health-report	Containers	Mem Used	Mem Avail	VCore Used
/default-rack		RUNNING	slave1:46290	slave1:8042	Fri Oct 28 12:48:37 -0700 2016		0	0 B	8 GB	0
/default-rack		RUNNING	slave2:37113	slave2:8042	Fri Oct 28 12:48:37 -0700 2016		0	0 B	8 GB	0

REFERENCE :KAUSHIK SHAKKARI
EMAIL : kaushik.shakkari@gmail.com

HADOOP MULTINODE INSTALLATION

master:50070/dfshealth.html#tab-overview		Search
Cluster ID:	CID-2d62c2e7-f156-4e8d-8fdd-c65e60d8eb97	
Block Pool ID:	BP-1517400136-192.168.220.180-1477676734424	

Summary

Security is off.

Safemode is off.

1 files and directories, 0 blocks = 1 total filesystem object(s).

Heap Memory used 78.89 MB of 159 MB Heap Memory. Max Heap Memory is 889 MB.

Non Heap Memory used 43.35 MB of 44.03 MB Committed Non Heap Memory. Max Non Heap Memory is -1 B.

Configured Capacity:	109.99 GB
DFS Used:	64 KB (0%)
Non DFS Used:	16.45 GB
DFS Remaining:	93.54 GB (85.04%)
Block Pool Used:	64 KB (0%)
DataNodes usages% (Min/Median/Max/stdDev):	0.00% / 0.00% / 0.00% / 0.00%
Live Nodes	2 (Decommissioned: 0)
Dead Nodes	0 (Decommissioned: 0)

HADOOP MULTINODE INSTALLATION

```
hduser@master: /usr/local/hadoop/logs
hduser@master:/usr/local/hadoop/logs$ hdfs dfsadmin -report
16/10/28 13:05:13 WARN util.NativeCodeLoader: Unable to load native code library for your platform... using builtin-java classes where applicable
Configured Capacity: 118101630976 (109.99 GB)
Present Capacity: 100438622208 (93.54 GB)
DFS Remaining: 100438556672 (93.54 GB)
DFS Used: 65536 (64 KB)
DFS Used%: 0.00%
Under replicated blocks: 0
Blocks with corrupt replicas: 0
Missing blocks: 0
Missing blocks (with replication factor 1): 0

-----
Live datanodes (2):

Name: 192.168.220.183:50010 (slave2)
Hostname: slave2
Decommission Status : Normal
Configured Capacity: 59050815488 (55.00 GB)
DFS Used: 32768 (32 KB)
Non DFS Used: 8831447040 (8.22 GB)
DFS Remaining: 50219335680 (46.77 GB)
DFS Used%: 0.00%
DFS Remaining%: 85.04%
Configured Cache Capacity: 0 (0 B)
Cache Used: 0 (0 B)
Cache Remaining: 0 (0 B)
Cache Used%: 100.00%
```